

FINM 32950: Intro to HPC in Finance

Lecture 2

Chanaka Liyanarachchi

March 31, 2025

Vectorization

Using Multicore Nodes

Vectorization

Vectorization in Action

- ▶ Last week: we introduced the idea of vectorization.
- ▶ Today: we look at how vectorization works using examples.
- ▶ We will look at 2 ways to introduce vectorization:
 1. Using intrinsics (low-level functions built in to the compiler)¹.
 2. Auto vectorization (using high-level language constructs).
- ▶ Intrinsics is a powerful technique but we do not use it directly in this course. We use intrinsics mainly to demonstrate/prove vectorization in action.
- ▶ In this course we prefer auto vectorization (for reasons we shall discuss later).

¹<https://docs.microsoft.com/en-us/cpp/intrinsics/compiler-intrinsics?view=vs-2019>

Representing Packed Data

- ▶ C++ defines data types for fundamental types.
- ▶ E.g.: int, double, float, short
- ▶ We cannot directly use them to store packed data.
- ▶ We need to use *vector extension types* to store packed data.
- ▶ The register size and fundamental data type determine how many data items we can store in a packed-register:
 - ▶ 128 bit register: can store 4 floats or 2 doubles
 - ▶ 256 bit register: can store 8 floats or 4 doubles
- ▶ Vector extension (packed data) types depend on:
 1. Size of the registers (register size depends on the architecture, SSE2: 128; AVX2: 256; AVX-512: 512)
 2. Fundamental data type (individual data item size depends on type, float: 32; double:64)

Packed Data Types

- ▶ For 128 bit registers:
 - ▶ `__m128`: stores 4 floats
 - ▶ `__m128d`: stores 2 doubles
 - ▶ `__m128i`: stores 4 ints
 - ▶ ...
- ▶ For 256 bit registers:
 - ▶ `__m256`: stores 8 floats
 - ▶ `__m256d`: stores 4 doubles
 - ▶ `__m256i`: stores 8 ints
 - ▶ ...
- ▶ There's a pattern:
 - ▶ prefix: `__m`
 - ▶ size of the register (e.g. 128, 256)
 - ▶ fundamental data type (i: int; d: double; default: float; ...)

Operations on Packed Data

- ▶ We must use intrinsic functions (packed instructions) to operate on packed data (i.e., we cannot use regular arithmetic operators (+, -, etc.) with packed types).
- ▶ Intrinsics provide a *family of functions* for a given operation (e.g. addition) for each packed type.
- ▶ Example: for SSE2:
 - ▶ `_mm128_add_epi32`: to add packed (4) integer values stored in a 128 bit packed register
 - ▶ `_mm128_add_epi16`: to add packed (8) short values stored in a 128 bit packed register
 - ▶ ...
- ▶ Example: for AVX2:
 - ▶ `_mm256_add_epi32`: to add packed (8) integer values stored in a 256 bit packed register
 - ▶ `_mm256_add_epi16`: to add packed (16) integer values stored in a 256 bit packed register
 - ▶ ...
- ▶ Ref: <https://software.intel.com/sites/landingpage/IntrinsicsGuide/#=undefined>

Intrinsics: Example 1

- ▶ Example addition for SSE processors:

```
#include <xmmintrin.h> // for SSE
#include <stdio.h> //for printf

void add_test()
{
    __m128 a = _mm_set_ps(1.0f, 2.0f, 3.0f, 4.0f);
    __m128 b = _mm_set_ps(5.0f, 6.0f, 7.0f, 8.0f);

    __m128 c = _mm_add_ps(a, b);

    // displaying results
    float* f = (float*)&c;
    printf("%f %f %f %f\n", f[0], f[1], f[2], f[3]);
}
```


Intrinsics: Example 2

- ▶ Let's look at another example to show the power of vectorization.
- ▶ Suppose we want to do a multiply and add operation on vectors:

```
for (int i=0; i<16; i++)  
{  
    d[i] = a[i] * b[i] + c[i];  
}
```

- ▶ This for loop uses 16 additions and 16 multiplications.
- ▶ We can do all of them in one operation:

```
__m512 a = /*packed data*/  
__m512 b = /*packed data*/  
__m512 c = /*packed data*/  
  
__m512 d = _mm512_fmadd_ps(a, b, c);
```

Vectorization using Intrinsics: Advantages and Disadvantages

- ▶ We briefly looked at intrinsics to show:
 1. How vectorization works
 2. Benefits/power of vectorization
- ▶ Intrinsics has some drawbacks:
 - ▶ Code is difficult to read, write, and maintain (try to write a Black Scholes pricer)
 - ▶ Code is not portable
- ▶ Next, we will look at how to achieve vectorization using high-level programming constructs.

Auto-Vectorization

Auto-Vectorization

- ▶ Modern compilers can vectorize code *automatically*. It is known as auto-vectorization.
 - ▶ The compiler generates packed SIMD instructions to replace loops.
- ▶ For auto-vectorization to work:
 - ▶ Code has to satisfy some requirements
 - ▶ Certain directives have to be used to encourage/force the compiler auto-vectorize.
- ▶ We need to understand how we can write code to maximize auto-vectorization:
 - ▶ If auto-vectorization occurred
 - ▶ What if it did not? we need to know why the compiler did not vectorize the loop/loops
 - ▶ So, we can do something to *encourage/force* compiler to auto-vectorize

Vectorization Using Intel C++ Compiler

- ▶ Modern compilers support auto vectorization, but details are vendor specific.
- ▶ We will limit our discussion to the **Intel C++ compiler** (works on Windows, Linux and Mac).
- ▶ General ideas about vectorization however are not vendor specific.

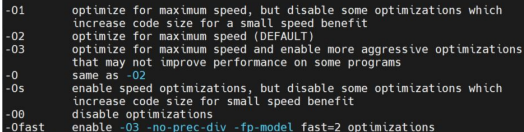
Auto-Vectorization: Example 1

- ▶ Let's add two vectors:

```
int main()
{
    const int N = 8;
    int a[N] = {1, 2, 3, 4, 5, 6, 7, 8};
    int b[N] = {1, 2, 3, 4, 5, 6, 7, 8};
    int c[N];

    for (int i = 0; i < N; ++i)
        c[i] = a[i] + b[i];
}
```

- ▶ Intel compiler supports following optimizations:

A screenshot of a terminal window showing various Intel compiler optimization flags and their descriptions. The text is as follows:

```
-O1      optimize for maximum speed, but disable some optimizations which
         increase code size for a small speed benefit
-O2      optimize for maximum speed (DEFAULT)
-O3      optimize for maximum speed and enable more aggressive optimizations
         that may not improve performance on some programs
-O       same as -O2
-Os      enable speed optimizations, but disable some optimizations which
         increase code size for small speed benefit
-O0      disable optimizations
-Ofast   enable -O3 -no-prec-div -fp-model fast=2 optimizations
```

- ▶ Use help for all available options:
`icc -help`

- ▶ Compiler is expected to auto-vectorize if optimization level 2 (O2) or higher is used.
- ▶ Will not vectorize if O1 or lower is used.
- ▶ Using O2:
`icc -O2 example1.cpp -o example1`
- ▶ Is the loop above auto vectorized?

Vectorization Report

- ▶ A *vectorization-report* shows what loops were vectorized and explains why any loop was not vectorized. To generate a vectorization report:
- ▶ On Linux:

`-qopt-report[=n]`

Generates an optimization report. Default destination is `<target>.optrpt`. Levels of 0 - 5 are valid.

Vectorization Report

- ▶ Each level provides different amount of information:
 - ▶ n=1: Loops successfully vectorized
 - ▶ n=2: Loops not vectorized; and the reason why not
 - ▶ n=3: Adds dependency information
 - ▶ n=4: Reports only non-vectorized loops
 - ▶ n=5: Reports only non-vectorized loops and adds dependency info
- ▶ The report file has the .REP extension on Windows and .optrpt on Linux.
- ▶ Additionally, we use the following to limit the output to vectorization only as the reports are generally lengthy and difficult to understand at first (you should experiment with and without this flag):
`-qopt-report-phase=vec`

- ▶ Let's create a vectorization report:
`icc -qopt-report=1 -O2 example1.cpp -o example1`
- ▶ In the report (example1.optrpt) we see that this loop is vectorized:

```
LOOP BEGIN at example1.cpp(9,2)
    remark #15300: LOOP WAS VECTORIZED
LOOP END
```

- ▶ Let's build the same program with optimizations disabled:
`icc -qopt-report=1 -O0 example1.cpp -o example1`
- ▶ The loop is not vectorized.
- ▶ Works as advertised.

Auto-Vectorization: Example 2

- Here's another example²:

```
int main()
{
    int sum = 0;
    for (int row = 0; row < ROWS; ++row)
    {
        for (int col = 0; col < COLS; ++col)
        {
            data[row][col] = row + col;
        }
    }
    for (int row = 0; row < ROWS; ++row)
    {
        for (int col = 0; col < COLS; ++col)
        {
            sum += data[row][col] + data[col][row];
        }
    }
}
```

²This example uses some arbitrary operations on a matrix, just for illustration.

- ▶ We have several (for) loops now.

```
icc -qopt-report=1 -O2 example2.cpp -o example2
```

- ▶ First inner loop is vectorized, but the second inner loop is not vectorized.

```
LOOP BEGIN at example2.cpp(13,7)
    remark #15300: LOOP WAS VECTORIZED
LOOP END
```

```
LOOP BEGIN at example2.cpp(21,7)
    remark #15335: loop was not vectorized:
LOOP END
```

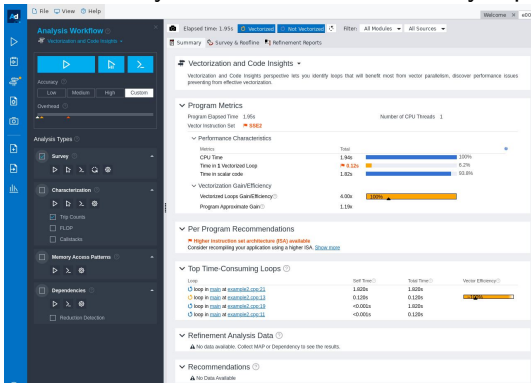
Intel Advisor

- ▶ Intel toolkit has many features to write performance sensitive code.
- ▶ Intel Advisor helps identify under optimized loops related info.
- ▶ Let's use Advisor to analyse Example 2 above.

Using Intel Advisor on Midway

- ▶ Load the modules:
`module use /software/intel/oneapi_hpc_2022.1/modulefiles`
`module load advisor/2022.0.0`
- ▶ Copy L2Demo.tar from /projects/finm32950/chanaka/ directory and untar (uncompress it).
- ▶ Change directory to L2Demo/advisor directory:
`cd L2Demo/advisor`
- ▶ Build:
`make`
- ▶ Collect survey data:
`advisor --collect=survey --project-dir=. ./example2`
- ▶ Generate survey report:
`advisor --report=survey --project-dir=.`
- ▶ Examine report:
`advisor-gui ./`
- ▶ Use "Show Results" to see the results.

► Use summary tab to view the summary report:



► Use survey and roofline tab to get details:

Analysis Workflow

Accuracy: Low Medium High Custom

Overhead: Low Medium High Custom

Analysis Types

- Survey
- Characterization
- Trips Counts
- FLDP
- Calculus
- Memory Access Patterns
- Dependencies
- Reduction Detection

Elapsed Time: 1.95s

Filter: All Modules All Sources Loops And Functions All Threads

Summary: Survey & Roofline Refinement Reports

Higher instruction set architecture (ISA) available
Consider recompiling your application using a higher ISA.

Function	Call Sites and Loops	Performance Issues	CPU Time	Total Time	Self Time	Type	Why No Vectorized?
at example2.cpp:13			1.820s	1.820s	0.000s	Scalar	vectorization
at example2.cpp:13			0.020s	0.020s	0.000s	Vectorized (Body)	
			1.840s	0.000s	0.000s	Function	
			1.840s	0.000s	0.000s	Function	

Source: example2.cpp:13 main

Line	Source	Total Time	%	Loop/Function Time
7	int main()			
8	{			
9	int sum = 0;			
10	for (int row = 0; row < ROWS; ++row)			
11	{			
12	loop in main at example2.cpp:11			
13	Scalar loop with instructions that use SIMD registers. Not vectorized; inner loop was not vectorized.			
14	No loop transformations applied.			
15	for (int col = 0; col < COLS; ++col)			
16	{			
17	loop in main at example2.cpp:13			
18	Vectorized body loop processor (int32 data type).			
19	Loop was unrolled by 4.			
20	{			
21	data[row][col] = row + col;			
22	}			
23	}			
24	for (int row = 0; row < ROWS; ++row)			
25	{			
26	loop in main at example2.cpp:11			
27	Scalar loop. Outer loop was not auto-vectorized; consider using SIMD directive.			
28	No loop transformations applied.			
29	{			
30	for (int col = 0; col < COLS; ++col)			
31	{			
32	loop in main at example2.cpp:13	0.010s		1.820s
33	Scalar loop. Not vectorized; vectorization possible but seems insufficient. Inner loop was unrolled by 3.			
34	}			
35	}			
36	}			
37	}			

Selected Total Time: 0s

Barriers to Vectorization

Example 3

- ▶ Let's look at this example:

```
#include <cstdlib>

int main()
{
    const int N = 8;
    float a[N], b[N], c[N];

    for (int i=0; i<N; ++i)
    {
        a[i] = rand() % 100;
        b[i] = rand() % 100;
    }

    for (int i = 0; i < N; ++i)
        c[i] = a[i] + b[i];
}
```

- ▶ We have 2 for loops. Will the compiler vectorize them?

- ▶ Let's compile this (O2 is default; we don't need to type it):

```
icc -qopt-report=1 example3.cpp
```

- ▶ Vectorization report:

```
LOOP BEGIN at example3.cpp(8,2)
    remark #25436: completely unrolled by 8
LOOP END
```

```
LOOP BEGIN at example3.cpp(14,2)
    remark #15300: LOOP WAS VECTORIZED
LOOP END
```

- ▶ Let's compile this again to find out why the first loop is not vectorized:

```
icc -qopt-report=2 example3.cpp
```

- ▶ Vectorization report:

```
LOOP BEGIN at example3.cpp(8,2)
```

```
    remark #15344: loop was not vectorized: vector  
    dependence prevents vectorization.
```

```
    First dependence is shown below. Use level 5 report  
    for details
```

```
LOOP BEGIN at example3.cpp(14,2)
```

```
    remark #15300: LOOP WAS VECTORIZED
```

```
LOOP END
```

- ▶ Now, the vectorization-report to guides us (and, compiling this again):

```
icc -qopt-report=5 example3.cpp
```

- ▶ Vectorization report:

```
LOOP BEGIN at example3.cpp(8,2)
    remark #15382: vectorization support: call to function
    rand() cannot be vectorized    [ example3.cpp(10,10) ]
```

- ▶ This loop is not vectorized because we are using rand() which doesn't support vectorization.

Example 4

- ▶ Another example of a loop that's not auto-vectorized:

```
for (int i = 0; i < N; ++i)
    if(i>3) c[i] = a[i] + b[i] + c[i-1];
```

- ▶ This loop won't be auto vectorized:

```
remark #15344: loop was not vectorized: vector dependence
prevents vectorization.
```

- ▶ Why?

- ▶ Let's unroll this loop to see what's going on:

```
c[0] = a[0] + b[0];  
c[1] = a[1] + b[1];  
.....  
c[4] = a[4] + b[4] + c[3];  
c[5] = a[5] + b[5] + c[4];  
.....
```

- ▶ We cannot calculate $c[4]$ and beyond in parallel due to dependency to other elements.
- ▶ This data dependence prevents vectorization.

Example 5

- ▶ Let's look at this example, which is similar to Black Scholes pricer where we are passing several arrays to a function:

```
void add(float *a, float *b, float *c,  
        float *d, float *e, int N)  
{  
    for (int i=0; i<n; ++i)  
        a[i] = b[i] + c[i] + d[i] + e[i];  
  
}
```

- ▶ Having many pointers may make it difficult for the compiler to figure out any dependence.
- ▶ Compiler may not know if the pointers are pointing to aliased memory.
- ▶ We have a data dependency problem if elements are overlapped.

- ▶ You may see a vectorization report similar to what's shown below, where one entry indicates the loop is vectorized:

```
LOOP BEGIN at example4.cpp(7,5)
<Multiversiomed v1>
    remark #25228: Loop multiversiomed for Data Dependence
    remark #15300: LOOP WAS VECTORIZED
LOOP END
```

- ▶ And, another entry indicates the same loop is not vectorized:

```
LOOP BEGIN at example4.cpp(7,5)
<Multiversiomed v2>
    remark #25436: completely unrolled by 8
LOOP END
```

- ▶ We have 2 versions for the same loop. What's going on?

Multiversions and Unrolling Loops

- ▶ Compiler may generate two versions (a vectorized version and a non-vectorized version) when it is not sure if vectorization is safe.
- ▶ Sometimes a compiler may unroll a loop instead of vectorizing it.
- ▶ Run time uses most appropriate one (i.e. vectorization is not guaranteed) so this is not ideal.

Example 6

- ▶ Loop count should remain constant during the execution of the loop for auto vectorization.

```
while (i<N)
{
    a[i] = a[i] * b[i];

    if (a[i] < 2) break;

    ++i;
}
```

remark #15520: loop was not vectorized: loop with multiple exits cannot be vectorized

- ▶ A loop must have a single entry and a single exit (entire loop must run) for auto vectorization.

Reading Vectorization Reports

- ▶ Examples above show some (not all) common vectorization reports.
- ▶ We can use them to learn how read and understand vectorization reports.
- ▶ You may see similar types of reports for other applications.

Improving Vectorization

- ▶ We saw some cases where auto vectorization is not used/possible.
- ▶ Shown below are some techniques to improve vectorization in such cases:
 1. Using vectorized implementations (e.g. SVMML)
 2. Compiler directives
 3. Writing code to support vectorization

1. Using vectorized implementations - Short Vector Math Library (SVML)

- ▶ A loop is not vectorized if a non vectorizable function is used in a loop.
- ▶ One solution is to use vectorized implementations of such functions.
- ▶ *Intel Short Vector Math Library* library provides vectorized implementations for some functions:

```
void test(float* a, float* b, int n)
{
    for (int i=0; i<n; ++i)
    {
        a[i] = sinf(b[i]);
    }
}
```

- ▶ A loop that uses vectorized functions will be auto vectorized.

- ▶ Support for vectorized implementations in the SVML include:
sin, cos, tan, asin, acos, atan, log, log2, log10, exp, exp2,
sinh, cosh, tanh, asinh, acosh, atanh, erf, erfc, erfinv, sqrt,
cbrt, trunc, round, ceil, floor, fabs, fmin, fmax, pow, atan2
- ▶ Reference:
<https://www.intel.com/content/www/us/en/docs/cpp-compiler/developer-guide-reference/2021-8/intrinsics-for-short-vector-math-library-ops.html>

2. Compiler Directives

- ▶ The compiler uses a static analysis to determine if vectorization safe (i.e produce correct results).
 - ▶ Won't vectorize any code unless it can guarantee correctness.
 - ▶ May ignore vectorization even when vectorization is possible due to ambiguity.
- ▶ Compilers provides directives so the programmer can guide/force vectorization when we (the programmer) know it is safe to vectorize such code.
- ▶ A guidance usually applies to a localized area (e.g. a function, a loop).

- ▶ The Intel C++ supports following pragma directives to help vectorization:
 - ▶ `simd`: instructs the compiler to enforce vectorization
 - ▶ `ivdep`: instructs the compiler to ignore assumed vector dependencies (guidance to the compiler)
 - ▶ `vector always`: force vectorization when "vectorization possible but seems inefficient"
- ▶ Related:
 - ▶ `novector`: tells the compiler not to vectorize a loop
 - ▶ `unroll`: unroll a loop
 - ▶ `nounroll`: don't unroll a loop
- ▶ Ref: <https://www.intel.com/content/www/us/en/docs/cpp-compiler/developer-guide-reference/2021-8/pragmas.html>
- ▶ We will discuss openmp directives next week.

Example 4 (cont..): simd Keyword

- Compiler generated two versions in this case, due to pointer aliasing (Example 4). We can force vectorization if we know pointer aliasing is not an issue in our application:

```
void add(float *a, float *b, float *c,  
         float *d, float *e, int n)  
{  
    #pragma simd  
    for (int i=0; i<n; i++)  
        a[i] = b[i] + c[i] + d[i] + e[i];  
}
```

Example 4 (cont..): ivdep Keyword

- Or, we could use:

```
void add(float *a, float *b, float *c,  
        float *d, float *e, int n)  
{  
    #pragma ivdep  
    for (int i=0; i<n; i++)  
        a[i] = b[i] + c[i] + d[i] + e[i];  
}
```

Example 2 (cont.): vector Keyword

- Compiler generated two versions for the second inner loop earlier (Example 2):

```
for (int row = 0; row < ROWS; ++row)
{
    for (int col = 0; col < COLS; ++col)
    {
        data[row][col] = row + col;
    }
}

for (int row = 0; row < ROWS; ++row)
{
    #pragma vector always
    for (int col = 0; col < COLS; ++col)
    {
        sum += data[row][col] + data[col][row];
    }
}
```

- Both inner loops are vectorized now.

3. Writing Code to Support Vectorization

► Following guidelines may improve auto vectorization:

1. Simplify loops. Avoid multiple exits and complex loop termination conditions.
2. Avoid data dependance in loops:

- E.g.: Move the first and/or the last iterations out of the loop to remove dependencies (known as loop peeling).

Before

```
for (int i=0; i<10; ++i)
{
    a[i] = b[i] + a[0];
}
```

After

```
a[0] = b[0] + a[0];
for (int i=1; i<10; ++i)
{
    a[i] = b[i] + a[0];
}
```

3. Avoid mixing types in the same loop.
4. ...

Unrolling Loops

- ▶ We cannot vectorize every loop. Unrolling can be helpful when vectorization is not possible.
- ▶ Loops is an important control structure that allow us to write concise and readable code.
- ▶ Unrolling is a technique used to optimize a loop by minimizing the cost of loop overhead:
 - ▶ Evaluate termination condition.
 - ▶ Update loop counter/counters.
- ▶ Unroll directives allow us to write code using loops but unroll them at compile time.

- ▶ Suppose we have a loop:

```
for (int i=0; i<4; ++i)
{
    a[i] = b[i] + c[i];
}
```

- ▶ Fully unrolling it:

```
a[0] = b[0] + c[0];
a[1] = b[1] + c[1];
a[2] = b[2] + c[2];
a[3] = b[3] + c[3];
```

- ▶ Unrolling by a factor of two:

```
for (int i=0; i<4; i+=2)
{
    a[i] = b[i] + c[i];
    a[i+1] = b[i+1] + c[i+1];
}
```

- ▶ Code is executed sequentially – not as efficient as vectorization.
- ▶ Performance improvements depend on the unrolling factor and the best factor can only be determined through measurements (discussed last week).

Vectorization: Usage


- ▶ Many popular libraries use vectorization to optimize operations:
 1. Eigen: http://eigen.tuxfamily.org/index.php?title=Main_Page
 2. RapidJSON: <https://github.com/Tencent/rapidjson>
 3. NumPy: https://www.pythonlikeyoumeanit.com/Module3_IntroducingNumpy/VectorizedOperations.html#

Remarks

- ▶ Vectorization is important.
- ▶ We discussed 2 ways to introduce vectorization:
 1. Using processor intrinsic functions
 2. Compiler auto vectorization
- ▶ Every application is different. There are no hard and fast rules that guarantee auto-vectorization.
- ▶ Recommendations/techniques described above may increase the success rate.

Using Multicore Nodes

- ▶ Let's go back to the report from vtune last time and focus on the *Effective CPU Utilization Histogram* section.



- ▶ We have over 90 CPU cores but we just use one.
- ▶ How do we use more than one CPU cores?

Parallel Programming - Introduction

Parallel Programming


- ▶ One obvious way get things done faster is to do more than one task simultaneously (in parallel):
 - ▶ This idea has been around for a long time.
 - ▶ It has become a very important topic now.
- ▶ We've seen vectorization parallelism already, and it works within one CPU core, giving us parallelism at the instruction level.
- ▶ Our next topic is parallelism using multithreading, allowing us to use many CPU cores.

Moore's Law

- ▶ Let's first try to understand why parallel programming using multi cores is needed.
- ▶ Moore's law says: every two years computing power doubles.
- ▶ Graphs below confirm that, but only upto about 2010. ³

50 Years of Technology Scaling

48 Years of Microprocessor Trend Data



³source:<https://www.datacenterknowledge.com/supercomputers/after-moore-s-law-how-will-we-know-how-much-faster-computers-can-go>

- ▶ Frequency has hit a limit due to hardware limitations (due to leakage currents in transistors).
- ▶ Single thread performance is still improving mainly due to vectorization.
- ▶ Trend now is to have more and more CPU cores.
 - ▶ Achieve same performance or better.
 - ▶ Less power consumption and heat generation.

Learning Parallel Programming

- ▶ New challenge for us (programmers):
 - ▶ We should know how to use multicores.
 - ▶ Otherwise, we are wasting a lot of computing power.
 - ▶ Learning how to write sequential code is not enough anymore.
- ▶ Parallel computing is not a highly-specialized/exotic topic anymore; it is something every programmer should know.

Parallel Computing: Example 1

- ▶ Suppose we want to add two arrays (vectors):

$$\begin{bmatrix} a_0 \\ a_1 \\ a_2 \end{bmatrix} + \begin{bmatrix} b_0 \\ b_1 \\ b_2 \end{bmatrix} = \begin{bmatrix} c_0 \\ c_1 \\ c_2 \end{bmatrix}$$

- ▶ We can write a simple for loop:

```
for (int i=0; i<N; ++i)
    c[i] = a[i] + b[i];
```

- ▶ This loop adds two values and assign it to $c[i]$ for each i in N .
- ▶ Usually when we think about programs (thus far), we think about doing things sequentially:
 - ▶ Use one data item from each array
 - ▶ Executing one instruction at a time
- ▶ Each addition is independent of the others (i.e. $c[0]$ doesn't depend on $c[1]$ and so on) – we can compute them in parallel.

Parallel Computing: Example 2

- ▶ Matrix multiplication is another example:

$$A_{I,K} B_{K,J} = C_{I,J} \quad c_{i,j} = \sum_{k=0}^{K-1} a_{i,k} b_{k,j} \quad \text{where } 0 \leq i < I \text{ and } 0 \leq j < J$$

$$\begin{pmatrix} a_{0,0} & a_{0,1} & \dots & a_{0,k-1} \\ a_{1,0} & a_{1,1} & \dots & a_{1,k-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{i-1,0} & a_{i-1,1} & \dots & a_{i-1,k-1} \end{pmatrix} \begin{pmatrix} b_{0,0} & b_{0,1} & \dots & b_{0,j-1} \\ b_{1,0} & b_{1,1} & \dots & b_{1,j-1} \\ \vdots & \vdots & \ddots & \vdots \\ b_{k-1,0} & b_{k-1,1} & \dots & b_{k-1,j-1} \end{pmatrix} \\ = \begin{pmatrix} c_{0,0} & c_{0,1} & \dots & c_{0,j-1} \\ c_{1,0} & c_{1,1} & \dots & c_{1,j-1} \\ \vdots & \vdots & \ddots & \vdots \\ c_{i-1,0} & c_{i-1,1} & \dots & c_{i-1,j-1} \end{pmatrix}$$

- ▶ Each element $c_{i,j}$ is independent of others – we can compute them in parallel.

Natural Parallelism

- ▶ Many applications in real life exhibit natural parallelism - vector addition and matrix multiplication are two examples.
- ▶ Our goal in this course is to identify and use natural parallelism to speed up program executions.
- ▶ We will explore 3 *technologies*:
 1. Vectorization ✓
 2. Multicore
 3. GPGPU
- ▶ FPGA is another interesting technology but we won't have time to discuss that topic.

Multithreading

Threads

- ▶ To use more than one core (multi cores), we break up a program into smaller *tasks*.
- ▶ We use *threads* to execute these tasks on different cores.
- ▶ This is generally known as multithreading.
- ▶ Our next topic is to explore how to use threads to do work in parallel.
- ▶ First, let's look at some important concepts about using threads.

Using Threads: Example 1

- ▶ Let's extend our familiar *Hello World* example to use a thread to perform a task.
- ▶ Task in this case is to write the greeting message to console.

```
#include <iostream>
#include <thread>

void Greeting()
{
    std::cout << "Hello, World" << std::endl;
}

int main()
{
    std::thread t(Greeting);

    t.join();
}
```

To build, on midway:

```
icc helloworld.cpp -o helloworld -lpthread
```

Hello World Example

- ▶ To create a thread, we use `std::thread`.
- ▶ The `std::thread` constructor takes a function as an argument.
- ▶ Thread execution begins in this function.
- ▶ Once the new thread is started, we have two threads in the program:
 1. main thread
 2. new thread (t)
- ▶ We use `join()` to make sure the calling thread (main) waits until new thread (t) finishes its execution.
- ▶ Thread is defined in `<thread>`
<http://www.cplusplus.com/reference/thread/thread/>

Thread Join and Detach

- ▶ Once a thread is created:
 - ▶ calling thread waits until the new thread completes its execution: use `join()`
 - ▶ calling thread continues execution without waiting for the new thread: use `detach()`

How Many Threads

- ▶ We can create a large number of threads in a program.
- ▶ How many threads can run in parallel depends on the number of cores/cpus.
- ▶ If we create a large number of threads, every thread may not run in parallel.
- ▶ The OS uses a scheduling algorithm to run the threads.

Detour: Lambdas (C++)

Lambdas

- ▶ We saw the `std::thread` constructor takes a function as an argument.
- ▶ There are other ways to pass an initial function to a thread:
 - ▶ using a lambda
 - ▶ using a function object

Lambda: Syntax

- ▶ The example below shows a very simple lambda which writes a message to console:

```
[] (string s)
{
    cout << s << endl;
}
```

- ▶ [] is called the capture clause; a lambda is always introduced by the capture clause.
- ▶ () parenthesis allows us to pass optional arguments to a lambda.
- ▶ {} lambda body is enclosed by the {} brackets.

Creating a Thread

- ▶ We can pass a lambda to a thread:

```
std::thread t([]()  
{  
    cout << "Hello, World" << endl;  
});  
  
t.join();
```


Lambda Syntax: Captures

- ▶ Lambdas allow us to *capture* parameters by value, or by reference.
- ▶ You can use *everything* notation to capture. In this case the compiler captures what's needed (used inside the body) by the lambda.
- ▶ [=] captures everything by value – allows reads but no write access.
- ▶ [&] captures everything by reference – allows read and write access.
- ▶ [=, &x] captures everything by value except x; x is captured by reference.
- ▶ [&, x] captures everything by reference except x; x is captured by value.

Using Threads: Example 2

► Matrix multiplication:

$$A_{I,K} B_{K,J} = C_{I,J} \quad c_{i,j} = \sum_{k=0}^{K-1} a_{i,k} b_{k,j} \quad \text{where } 0 \leq i < I \text{ and } 0 \leq j < J$$

$$\begin{pmatrix} a_{0,0} & a_{0,1} & \dots & a_{0,k-1} \\ a_{1,0} & a_{1,1} & \dots & a_{1,k-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{i-1,0} & a_{i-1,1} & \dots & a_{i-1,k-1} \end{pmatrix} \begin{pmatrix} b_{0,0} & b_{0,1} & \dots & b_{0,j-1} \\ b_{1,0} & b_{1,1} & \dots & b_{1,j-1} \\ \vdots & \vdots & \ddots & \vdots \\ b_{k-1,0} & b_{k-1,1} & \dots & b_{k-1,j-1} \end{pmatrix} \\ = \begin{pmatrix} c_{0,0} & c_{0,1} & \dots & c_{0,j-1} \\ c_{1,0} & c_{1,1} & \dots & c_{1,j-1} \\ \vdots & \vdots & \ddots & \vdots \\ c_{i-1,0} & c_{i-1,1} & \dots & c_{i-1,j-1} \end{pmatrix}$$

- ▶ We can write a serial program:

```
void matrix_multiply(const matrix& m1, const matrix& m2,
    matrix& m3, int rows, int columns)
{
    for (int i = 0; i < rows; ++i)
    {
        for (int j = 0; j < columns; ++j)
        {
            m3[i][j] = 0;
            for (int k = 0; k < rows; ++k)
            {
                m3[i][j] += m1[i][k]*m2[k][j];
            }
        }
    }
}
```

- ▶ Inputs (m1 and m2) are passed by const reference.
- ▶ Output (m3) passed by reference.

- ▶ This is a naturally parallel program - each $c_{i,j}$ can be computed in parallel.
- ▶ We can parallelize it in many different ways:
 - ▶ one thread to calculate each element $c_{i,j}$ – we need $\text{numRow} * \text{numColumn}$ number of threads
 - ▶ one thread for one row – we need numRow number of threads
 - ▶ one thread for one column – we need numColumn number of threads
- ▶ Which choice is better?

► Suppose we use one thread for each element:

```
void multiply_matrix(const matrix& m1, const matrix& m2,
                    matrix& m3, int rows, int columns)
{
    vector<thread> threads(rows*columns);

    for (int i = 0; i < rows; ++i)
    {
        for (int j = 0; j < columns; ++j)
        {
            int idx = i*columns + j;
            threads[idx] = thread(&m3[i][j], m1, m2);
            {
                m3[i][j] = 0;
                for (int k = 0; k < columns; ++k)
                {
                    m3[i][j] += m1[i][k] * m2[k][j];
                }
            }
        }
    }

    for (thread& t : threads)
    {
        t.join();
    }
}
```

- Or, we can to use one thread for one row:


```
void multiply_matrix(const matrix& m1, const matrix& m2,
    matrix& m3, int rows, int columns)
{
    vector<thread> threads(rows);

    for (int i = 0; i < rows; ++i)
    {
        threads[i] = thread([=,&m3]()
        {
            for (int j = 0; j < columns; ++j)
            {
                m3[i][j] = 0;
                for (int k = 0; k < rows; ++k)
                {
                    m3[i][j] += m1[i][k] * m2[k][j];
                }
            }
        });
    }

    for (thread& t : threads)
    {
        t.join();
    }
}
```

- ▶ Note: This program does not use any special features of Intel compiler (i.e. you may use Visual C++ or any other compiler).
- ▶ We have 3 implementations now:
 1. sequential work – one thread (i.e. main thread) computes all elements
 2. one thread computes one element
 3. one thread computes one row
- ▶ Which program/programs will run faster?

- ▶ Let's run these programs and measure time.
- ▶ Results depend on the system and the matrix sizes etc.
- ▶ Shown below is what I saw for one case (see demo for code) on my laptop:



```
Microsoft Visual Studio Debug Console  
(Serial) Time elapsed 1398 ms  
(Thread per element) Time elapsed 70039 ms  
(Thread per row) Time elapsed 571 ms
```


Observations

- ▶ Threads allow us to gain performance.
- ▶ More threads doesn't necessarily mean better performance.
Why?
- ▶ When we create more threads than what we can run in parallel, we create overhead.
- ▶ When we use threads we have to pay attention to the overhead.

Sharing Data

- ▶ All threads share same data.
- ▶ Multiple threads can read data at the same time (e.g., matrix multiplication above).
- ▶ Can more than one thread write to the same data at the the same time?
- ▶ Can one thread write while another one reads the same data?
- ▶ To answer these questions, let's look at an example.

Threads: Example 3

- Let's look at the code snippet below:

```
unsigned long counter = 0;

int numThreads = 10;
vector<thread> threads(numThreads);

for (int j = 0; j < numThreads; ++j)
{
    threads[j] = thread([&counter]()
    {
        for (int i = 0; i < 100000; ++i)
        {
            counter++;
        }
    });
}

for (auto& t : threads) t.join();

cout << counter;
```

► To build and run (on Midway):

1. Get a compute node with 8 CPU cores

```
sinteractive --time=0:30:0 --ntasks=8 --account=finm32950
```

2. Build:

```
icc counter.cpp -o counter -lpthread
```

3. Run:

```
./counter
```

► What's the value of counter?

Counter Example: Possible Scenario 1

- ▶ Suppose we have two threads t1 and t2 running in parallel:
 - ▶ t1 reads the value of counter, counter is 0
 - ▶ t2 reads the value of counter, counter is 0
 - ▶ t1 increments it, counter is 1
 - ▶ t2 increments it, counter is 1
 - ▶ after 2 increments, the value of counter is 1
- ▶ When we use threads, we should expect threads to interleave read and modify operations as shown below:

t1	t2
-----	-----
registerA = counter	registerB = counter
counter = registerA + 1	counter = registerB + 1

v (time/clock-cycle)

- ▶ In this case, the counter will be incremented only once after two add operations.

Counter Example: Possible Scenario 2

- ▶ Suppose we have two threads t1 and t2 running in parallel:
 - ▶ t1 reads the value of counter, counter is 0
 - ▶ t1 increments it, counter is 1
 - ▶ t2 reads the value of counter, counter is 1
 - ▶ t2 increments it, counter is 2
 - ▶ after 2 increments the value of counter is 2

t1

registerA = counter

counter = registerA + 1

t2

registerB = counter

counter = registerB + 1

- ▶ In this case, the counter will be incremented twice (correctly) after two add operations.

Race Conditions

- ▶ If two (or more) threads write to the same shared data, the result depends on the order in which the threads ran.
- ▶ This problem (bug) is known as a *race condition*.
- ▶ How do we avoid race conditions?

Critical Region

- ▶ The counter is a shared variable in our example.
- ▶ Incrementing (modifying) it by two or more threads at the same time leads to a race condition.
- ▶ The part where shared resource is accessed is known as the critical section/region

```
for (int j = 0; j < numThreads; ++j)
{
    threads[j] = thread([&counter]()
    {
        for (int i = 0; i < 100000; ++i)
        {
            counter++;
        }
    });
}
```

- ▶ Only one thread *should* execute the code in the critical section at a time – known as mutual exclusion.

Using a Lock

- ▶ One way to achieve mutual exclusion is to use a lock to protect shared data.
- ▶ We can use a lock known as mutex (*mutual exclusion*).
- ▶ To protect shared data using a mutex involves:
 - ▶ lock the mutex before shared data is accessed
 - ▶ unlock the mutex after shared data is accessed
- ▶ Only one thread can lock the mutex at a time.
- ▶ Everyone else has to wait until the thread holding the mutex releases it.
- ▶ <http://en.cppreference.com/w/cpp/thread/mutex>

- ▶ We can re-write the counter example using a mutex:

```
for (int j = 0; j < numThreads; ++j)
{
    threads[j] = thread([&]()
    {
        for (int i = 0; i<100000; ++i)
        {
            count_mutex.lock();
            counter++;
            count_mutex.unlock();
        }
    });
}
```

- ▶ We are guarding the critical region using a mutex.
- ▶ Race condition is avoided.

Deadlocks

- ▶ What happens if we lock a mutex and forget to unlock it?
- ▶ No thread will be able to access the shared resource again.
- ▶ This will lead to a problem, known as a *deadlock*.
- ▶ Deadlocks can happen due to exceptions:

```
mutex.lock();  
  
critical_region; //can throw an exception  
  
mutex.unlock();
```

std::lock_guard

- ▶ The `std::lock_guard` class implements the RAI technique⁴ for a mutex:
 - ▶ locks the mutex when the `lock_guard` object is constructed (in the constructor)
 - ▶ unlocks the mutex when the `lock_guard` object is destroyed (in the destructor)
- ▶ `std::lock_guard`:
 - ▶ prevents errors due to forgetting to release a lock
 - ▶ guarantees exception safety
- ▶ Defined in `<mutex>`

⁴discussed in winter course

- ▶ The code snippet below shows how to lock a piece of shared data using a `std::mutex` using this technique.

```
unsigned long count = 0;
std::mutex count_mutex;

int numThreads = 10;
vector<thread> threads(numThreads);

for (int j = 0; j < numThreads; ++j)
{
    threads[j] = thread([&]()
    {
        for (int i = 0; i<100000; ++i)
        {
            std::lock_guard<std::mutex> guard(count_mutex);
            count++;
        }
    });
}
```

Atomics

- ▶ Another way to handle critical regions is to use atomic operations.
- ▶ Atomic operations are performed as one (all or none) operation so the other threads see the operation as a single operation.
- ▶ Atomic operations do not require us to use locks – no need to worry about locking related issues.
- ▶ Number of atomic operations are limited – to handle complicated critical regions we still need to use locks.
- ▶ Atomic types are defined in atomic header (<http://en.cppreference.com/w/cpp/header/atomic>).

- We can write the counter example using atomics as follows:

```
atomic<unsigned long> counter;  
counter = 0;  
  
for (int j = 0; j < numThreads; ++j)  
{  
    threads[j] = thread([&]()  
    {  
        for (int i = 0; i<100000; ++i)  
        {  
            counter++;  
        }  
    }));  
}
```